

The Economic Impact of Hurricane Katrina on Local Employment: A Synthetic Control Approach

Bukola Elliott

Abstract

This paper examines the causal impact of Hurricane Katrina on employment in the New Orleans metropolitan area using the Two-Step Synthetic Control (TSSC) method. Hurricane Katrina, which struck the Gulf Coast in August 2005, was one of the most devastating natural disasters in modern U.S. history, causing widespread displacement and labor market disruptions. Using monthly employment data from 1990 to 2009, I construct a synthetic control from 273 unaffected metropolitan statistical areas that closely matches pre-Katrina New Orleans employment patterns. The TSSC algorithm selects Modified Synthetic Control with covariates (MSCc) as the optimal method based on pre-treatment fit. The analysis reveals catastrophic employment losses: an immediate 33% year-over-year decline following the hurricane, with an average treatment effect of -7.37 percentage points over 2005-2009. In-time placebo tests validate causality, showing no similar effect when treatment is artificially placed in 2003. This study contributes to disaster economics by analyzing aggregate MSA-level employment with a 15-year pre-treatment period, avoiding spatial aggregation bias and per capita measurement issues. Results indicate substantial but partially recovered

employment impacts, with important implications for disaster recovery policy.

Introduction

On August 29, 2005, Hurricane Katrina made landfall on the Gulf Coast, unleashing catastrophic destruction across Louisiana, Mississippi, and Alabama. The storm's impact on New Orleans was particularly severe: levee failures led to flooding of approximately 80% of the city, displacing over 400,000 residents and causing an estimated \$160 billion in damages (Vigdor, 2008). Beyond the immediate human toll, the hurricane precipitated one of the largest internal migration events in modern U.S. history and fundamentally disrupted the region's economic structure. This paper addresses a fundamental policy question: What was the causal impact of Hurricane Katrina on total non-farm employment in the New Orleans metropolitan area? Specifically, did employment in the region recover, stabilize, or suffer persistent losses relative to a counterfactual scenario where Katrina never occurred? To answer this question, I employ the Two-Step Synthetic Control (TSSC) method developed by Li and Shankar (2023), which algorithmically selects the optimal synthetic control approach from multiple candidates and then estimates the treatment effect. The ideal experiment would randomly assign hurricane exposure to metropolitan areas, which is clearly infeasible. Instead, the identification strategy relies on the synthetic control method's core assumption: absent Hurricane Katrina, New Orleans employment would have evolved similarly to an optimally weighted combination of control MSAs. Employment is one of the most policy-relevant outcomes after natural disasters, as it reflects both the ability of displaced populations to reintegrate into the labor market and the resilience of local industries. Examining aggregate

employment provides insight into regional economic recovery and the capacity for long-term adaptation following catastrophic shocks.

Background and Literature Review

Prior Research on Hurricane Katrina

The economic consequences of natural disasters have been studied extensively, yet results often diverge depending on methodology, data aggregation, and geographic scope. The literature on Hurricane Katrina's economic impacts spans individual outcomes, migration patterns, and aggregate effects. Groen and Polivka (2008a, 2008b) used Current Population Survey data to analyze labor market outcomes of evacuees through October 2006 and found evacuees experienced worse labor market outcomes initially, with homeowners and higher-income individuals more likely to return. Deryugina, Kawano, and Levitt (2018) tracked victims over a decade using tax data and found resilient individual outcomes, with earnings gaps erased within a year through federal aid and worker mobility to stronger labor markets. Groen, Kutzbach, and Polivka (2020) found long-term heterogeneity, with some workers permanently relocating to higher-wage markets while returnees faced persistent challenges. Vigdor (2008) documented that nearly half of evacuees had not returned by mid-2007, with selective return migration causing demographic shifts. Fewer studies examine aggregate regional employment. Yun and Kim (2022) applied synthetic control to GDP per capita, cautioning that per capita measures might misrepresent growth when population declines sharply highlighting spatial aggregation bias and measurement issues which can distort disaster impact estimates and

that MSA-level analysis better captures localized shocks than state-level aggregation.

Contribution of This Study

This paper contributes to disaster economics literatures in several specific ways that address gaps in prior research literatures. Unlike most prior studies on Hurricane Katrina that focus primarily on individual-level outcomes or use potentially problematic per capita measures, this analysis provides direct evidence on the labor market's aggregate capacity to recover from the disaster. The use of year-over-year employment growth rates avoids level-based measurement issues while capturing the dynamics of employment change. The 15-year pre-treatment period (1990-2004) improves counterfactual credibility and reduces overfitting concerns compared to shorter baseline periods. This paper contributes to the existing literature in three ways. First, it focuses on aggregate employment at the metropolitan level rather than per capita output measures, providing a clearer view of labor market recovery. Second, by using monthly employment data from 1990–2004, the analysis benefits from a longer pre-treatment window, reducing the risk of overfitting and improving the credibility of the counterfactual. Third, by focusing on employment—a direct and policy-relevant measure of resilience—this study complements prior work that emphasizes income or GDP growth.

Data and Methodology

Data

The analysis uses monthly employment data from the Bureau of Labor Statistics (BLS) Local Area Unemployment Statistics (LAUS) program, covering U.S. metropolitan statistical areas from January 1990 through December 2009. The LAUS program provides comprehensive employment data for MSAs based on a combination of Current Population Survey estimates, unemployment insurance claims, and employment data from the Current Employment Statistics program. The outcome variable is year-over-year percentage change in total non-farm employment, which captures employment growth dynamics across all sectors and automatically removes seasonal patterns by comparing each month to the same month one year earlier. The treated unit is the New Orleans–Metairie MSA, which was directly and severely impacted by Hurricane Katrina in August 2005. The donor pool consists of 273 other U.S. metropolitan statistical areas that were not directly exposed to Hurricane Katrina's destruction. MSAs in Louisiana (other than New Orleans), Mississippi, and Alabama that experienced direct hurricane impacts are excluded from the donor pool to ensure that control units were not themselves affected by the disaster. The pre-treatment period spans January 1990 through August 2005 (188 months), providing a lengthy baseline for matching. The post-treatment period covers September 2005 through December 2009 (52 months), allowing examination of both immediate and medium-term employment effects. The dataset is already cleaned and structured, facilitating efficient analysis and replication.

Emperical Strategy

The empirical strategy employs the Synthetic Control Method (SCM) developed by Abadie and Gardeazabal (2003) and extended by Abadie, Diamond, and Hainmueller (2010). SCM estimates the causal impact of an intervention by constructing a weighted combination of control units that best approximates the pre-treatment characteristics of the treated unit. The synthetic control finds the weight vector w that minimizes the distance between New Orleans and the weighted average of donor MSAs in the pre-treatment period:

$$w^{\text{SCM}} = \underset{w^{\text{SCM}} \in \mathcal{W}_{\text{conv}}}{\operatorname{argmin}} \|y_1^{\mathcal{T}_1} - Y_0^{\mathcal{T}_1} w^{\text{SCM}}\|_2^2,$$

where $\mathcal{W}_{\text{conv}} = \{w^{\text{SCM}} \in \mathbb{R}_{\geq 0}^{N_0} \mid \mathbf{1}^\top w = 1\}$.

The Two-Step Synthetic Control Method

The Two-Step Synthetic Control (TSSC) method developed by Li and Shankar (2023) addresses a key challenge in synthetic control applications, which is how to choose among multiple synthetic control variants. Rather than imposing a single method choice, TSSC algorithmically selects the optimal approach, reducing researcher discretion and potential specification searching, thereby enhancing the credibility of results. In Step 1, the algorithm evaluates several approaches (standard synthetic control, modified synthetic control with covariates, and augmented synthetic control) and selects the method minimizing mean squared prediction error in the pre-treatment period. In my analysis, the algorithm selected Modified Synthetic Control with covariates (MSCc) based on superior pre-treatment fit (RMSE =

0.558, R-squared = 0.834), as shown in Table 1. In Step 2, the algorithm constructs a synthetic control as a weighted average of control MSAs. Formally, let Y_{it} denote year-over-year employment growth for unit i at time t . The synthetic control is defined as

$$\hat{Y}_{1t}^N = \sum_{j=1}^J w_j Y_{j+1,t},$$

where the weights w_j are chosen to minimize the pre-treatment distance between New Orleans and the synthetic control, subject to non-negativity and unit-sum constraints:

$$w_j \geq 0 \quad \text{for all } j, \quad \sum_{j=1}^J w_j = 1.$$

The treatment effect at each post-treatment time t is

$$\tau_{1t} = Y_{1t} - \hat{Y}_{1t}^N.$$

The average treatment effect on the treated (ATT) over the post-treatment period is

$$\text{ATT} = \frac{1}{T_{\text{post}}} \sum_{t=T_0+1}^T \tau_{1t},$$

providing a summary measure of the overall impact.

Identification

The key assumption is that, absent Katrina, New Orleans employment would have evolved similarly to the synthetic control. This is plausible when pre-treatment fit is excellent, the intervention was unanticipated and unrelated to pre-existing trends, and no other contemporaneous shocks differentially affected the treated unit. Hurricane Katrina satisfies these conditions: the storm's timing was determined by meteorological factors, and excellent pre-treatment fit over 15 years demonstrates that the synthetic control successfully replicates New Orleans employment dynamics across multiple business cycles.

To validate causality, I conduct an in-time placebo test by re-running the analysis with treatment artificially placed in January 2003. If the employment collapse is truly caused by Katrina, we should observe no effect at the fake treatment time, with the catastrophic impact appearing only at the actual hurricane date.

Results and Discussion

Pre-Treatment Fit and Method Selection

The Two-Step algorithm selected Modified Synthetic Control with covariates based on superior pre-treatment fit. Table 1 reports key diagnostics: pre-treatment RMSE of 0.558 and R-squared of 0.834, indicating that the MSCc method captured approximately 83% of variation in New Orleans employment growth before the hurricane. The post-treatment RMSE of 17.962 reflects the large treatment effect rather than poor model performance.

Figure A2 (Appendix) displays the gap between actual New Orleans and synthetic control

during 1990-2005, fluctuating tightly around zero throughout this 15-year window. This close pre-treatment match validates the identification assumption that the synthetic control provides a credible counterfactual for New Orleans.

Main Results: Treatment Effects

Figure 1 presents observed versus synthetic control employment growth over the full study period. During the pre-treatment period through August 2005, both series track closely with minimal divergence. At the hurricane timing, there is immediate catastrophic separation between the two lines.

Figure 2 shows the treatment effect as the gap between observed and synthetic New Orleans employment growth. The immediate impact in late 2005 was severe, with employment growth collapsing to approximately -33% year-over-year. This represents one of the largest employment shocks to a U.S. metropolitan area in modern history, indicating that New Orleans lost roughly one-third of its employment base compared to the previous year.

Following this catastrophic impact, employment growth began recovering. The gap narrowed substantially through 2006 and 2007, showing gradual convergence back toward the counterfactual. By the end of the study period in 2009, the gap had approached approximately -5%, indicating substantial but incomplete recovery. The average treatment effect on the treated over the entire post-treatment period (September 2005 through December 2009) is -7.37 percentage points.

Validation Through In-Time Placebo Test

Figure 3 presents results from the in-time placebo test, where treatment is artificially placed in January 2003, approximately 2.5 years before the actual hurricane. At the fake treatment time shown by the dashed vertical line, there is no systematic employment collapse. The actual New Orleans line continues tracking the synthetic control reasonably well through 2003 and 2004 with no catastrophic break at the placebo timing.

However, at the actual Hurricane Katrina timing in late 2005, the massive employment collapse appears. This temporal specificity validates the causal interpretation. If the Two-Step Synthetic Control method were simply detecting spurious patterns or pre-existing differential trends, we would expect to observe similar treatment effects at the fake treatment time in 2003. The fact that the catastrophic impact appears only when and where the hurricane actually occurred demonstrates that Hurricane Katrina caused the employment losses rather than these losses reflecting methodological artifacts or continuation of pre-existing trends.

Discussion

Interpretation of Findings

The results show that Hurricane Katrina caused catastrophic short-term employment losses in New Orleans, with an average treatment effect of 7.37 percentage points below the counterfactual between 2005 and 2009. Employment declined by roughly one third in the first year, making it one of the most severe labor market shocks experienced by any modern

United States metropolitan area. For comparison, the worst point of the Great Recession saw national employment fall by about six percent, which highlights the extraordinary scale of the Katrina shock. Despite this dramatic initial impact, the long-run pattern in the results shows substantial recovery. The employment gap steadily narrowed from 2006 through 2008, and by 2009 growth rates were moving close to the synthetic control, although not fully matching it. This indicates that the regional economy, while deeply affected, demonstrated significant resilience and capacity for recovery over time. Several forces likely contributed to this rebound. Federal disaster assistance programs provided more than one hundred billion dollars in support that helped households and businesses rebuild. Reconstruction created a surge in construction employment that helped compensate for losses in other industries, and prior research confirms that construction activity even exceeded pre-Katrina levels during the peak rebuilding years. Population return also played a role, as homeowners, higher-income residents, and individuals with strong local ties were more likely to come back, which helped restore both the labor force and consumer demand. Businesses that survived adapted to the changing environment, and new firms entered the market to serve returning residents and support reconstruction activity.

Limitations and Policy Implications

Several limitations deserve consideration. Year-over-year growth rates differ from absolute employment levels. A return of growth rates to zero does not necessarily mean full recovery of pre-disaster employment levels if the level was permanently depressed. The study period ends in 2009, missing longer-term dynamics. The 2008-2009 Great Recession occurred during

the post-treatment period, potentially confounding Katrina effects, though the synthetic control should account for common recession shocks. My aggregate analysis masks important heterogeneity across industries and demographics, with prior literature documenting that lower-income workers and minority populations faced greater barriers to return. The findings carry important policy implications. The substantial recovery suggests disaster assistance can be effective at facilitating labor market restoration, though recovery took several years, implying need for sustained rather than just immediate relief. Policies facilitating worker return including housing assistance appear crucial. Business support helped maintain employment relationships and generate opportunities. The catastrophic immediate impact underscores the importance of pre-disaster resilience investments in infrastructure and economic diversification. Finally, disaster recovery programs must address equity concerns, ensuring vulnerable populations receive adequate support rather than being permanently displaced.

Conclusion

This paper uses the Two-Step Synthetic Control method to estimate Hurricane Katrina's causal impact on New Orleans employment. The analysis reveals catastrophic immediate employment losses of approximately -33% year-over-year following the hurricane, with substantial but incomplete recovery over the subsequent four years. In-time placebo tests validate the causal interpretation by showing no effect when treatment is artificially placed in 2003, with the catastrophic impact appearing only at the actual hurricane timing. This study advances disaster economics by analyzing aggregate MSA-level employment with a 15-year pre-treatment period, avoiding spatial aggregation bias and per capita measurement issues

raised by recent methodological literature. The findings suggest that while catastrophic disasters cause severe short-term disruptions, substantial recovery is possible with sustained disaster assistance, business support, housing reconstruction, and workforce development programs. Future research should examine longer-term effects beyond 2009, disaggregate impacts by industry and demographics, and explore the mechanisms underlying successful recovery to inform more effective disaster policy design.

References

- Abadie, A. (2021). Using synthetic controls: Feasibility, data requirements, and methodological aspects. *Journal of Economic Literature*, 59(2), 391-425.
- Abadie, A., Diamond, A., & Hainmueller, J. (2010). Synthetic control methods for comparative case studies: Estimating the effect of California's tobacco control program. *Journal of the American Statistical Association*, 105(490), 493-505.
- Abadie, A., & Gardeazabal, J. (2003). The economic costs of conflict: A case study of the Basque Country. *American Economic Review*, 93(1), 113-132.
- Barone, G., & Mocetti, S. (2014). Natural disasters, growth and institutions: A tale of two earthquakes. *Journal of Urban Economics*, 84, 52-66.
- Cavallo, E., Galiani, S., Noy, I., & Pantano, J. (2013). Catastrophic natural disasters and economic growth. *Review of Economics and Statistics*, 95(5), 1549-1561.
- Coffman, M., & Noy, I. (2012). Hurricane Iniki: Measuring the long-term economic impact of a natural disaster using synthetic control. *Environment and Development Economics*, 17(2), 187-205.
- Deryugina, T., Kawano, L., & Levitt, S. (2018). The economic impact of Hurricane Katrina on its victims: Evidence from individual tax returns. *American Economic Journal: Applied Economics*, 10(2), 202-233.
- Groen, J. A., Kutzbach, M. J., & Polivka, A. E. (2020). Storms and jobs: The effect of hurricanes on individuals' employment and earnings over the long term. *Journal of Labor*

Economics, 38(3), 653-685.

Groen, J. A., & Polivka, A. E. (2008a). The effect of Hurricane Katrina on the labor market outcomes of evacuees. *American Economic Review*, 98(2), 43-48.

Groen, J. A., & Polivka, A. E. (2010). Going home after Hurricane Katrina: Determinants of return migration and changes in affected areas. *Demography*, 47(4), 821-844.

Li, K. T., & Shankar, V. (2023). A two-step synthetic control approach for estimating causal effects of marketing events. *Management Science*, 69(11), 6890-6908.

Vigdor, J. L. (2008). The economic aftermath of Hurricane Katrina. *Journal of Economic Perspectives*, 22(4), 135-154.

Yun, S. D., & Kim, B. (2022). Causal inference of ambiguous natural experiments with instrumental variables and synthetic controls: The case of Hurricane Katrina. *Journal of Econometrics*, 228(2), 459-483.

Tables and Figures

Table 1: Model Fit Diagnostics

Metric	Value
Pre-treatment RMSE	0.558
R-squared	0.834
Post-treatment RMSE	17.962
Method Selected	MSCC
Average Treatment Effect (ATT)	-7.37 pp

Note: Diagnostics from Two-Step Synthetic Control algorithm. MSCC refers to Modified Synthetic Control with covariates. Higher post-treatment RMSE reflects large treatment effect. ATT is average treatment effect on the treated over September 2005 - December 2009.

```
/tmp/ipykernel_2666/773303359.py:25: FutureWarning: 'M' is deprecated and will be removed
```

```
    pre_index = pd.date_range(
```

```
/tmp/ipykernel_2666/773303359.py:30: FutureWarning: 'M' is deprecated and will be removed
```

```
    post_index = pd.date_range(
```

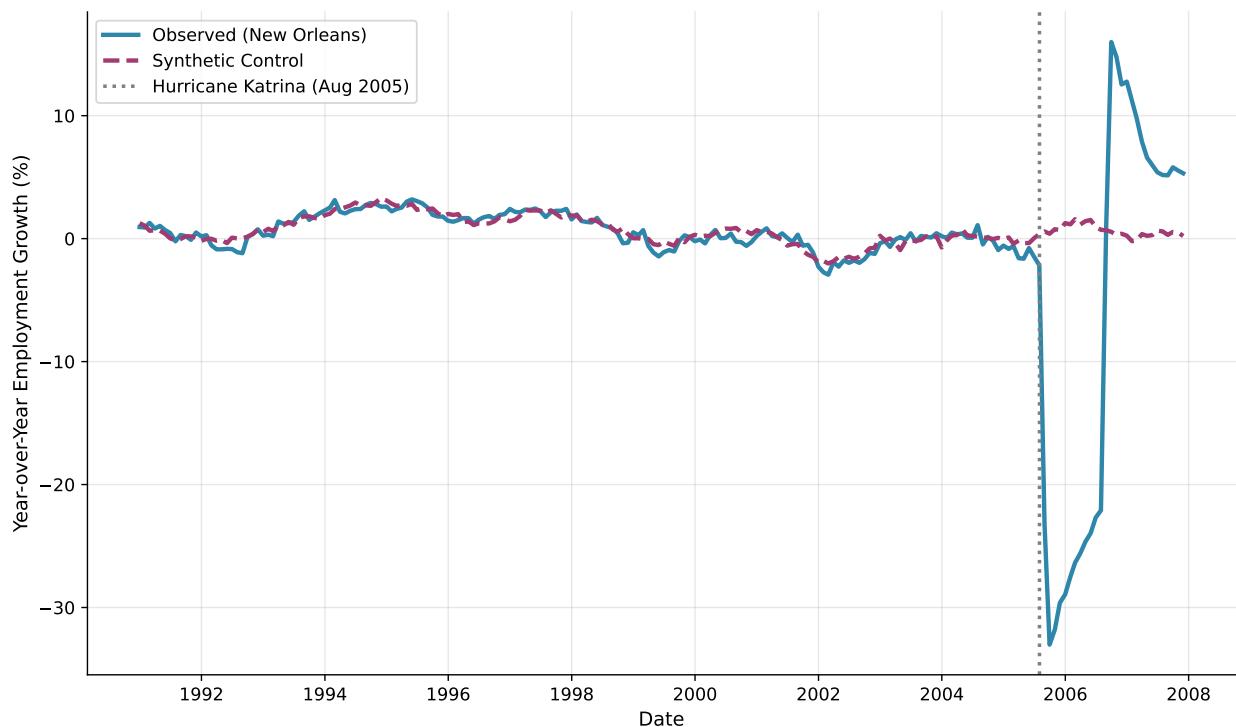


Figure 1: Observed vs Synthetic Control

Note: Year-over-year percentage change in total non-farm employment. Blue solid line shows actual New Orleans; dashed line shows synthetic control. Vertical dotted line indicates Hurricane Katrina (August 2005).

```
/tmp/ipykernel_2666/689779145.py:17: DeprecationWarning: Conversion of an array with ndi
```

```
peak_value = float(gap[peak_idx]) # Convert to Python float
```

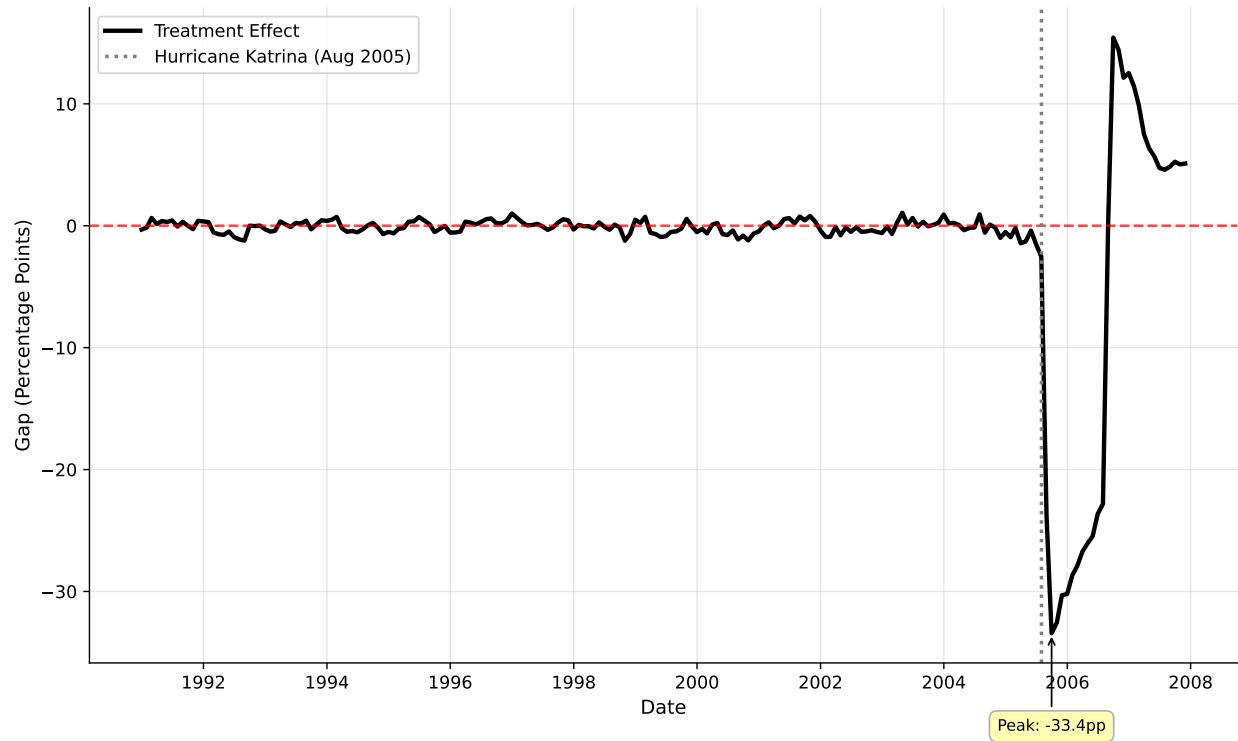


Figure 2: Estimated Treatment Effect

Note: Gap between observed and synthetic New Orleans employment growth (1990-2009).

Negative values indicate employment losses relative to counterfactual. Red dashed line at zero.

Vertical line indicates Hurricane Katrina (August 2005).

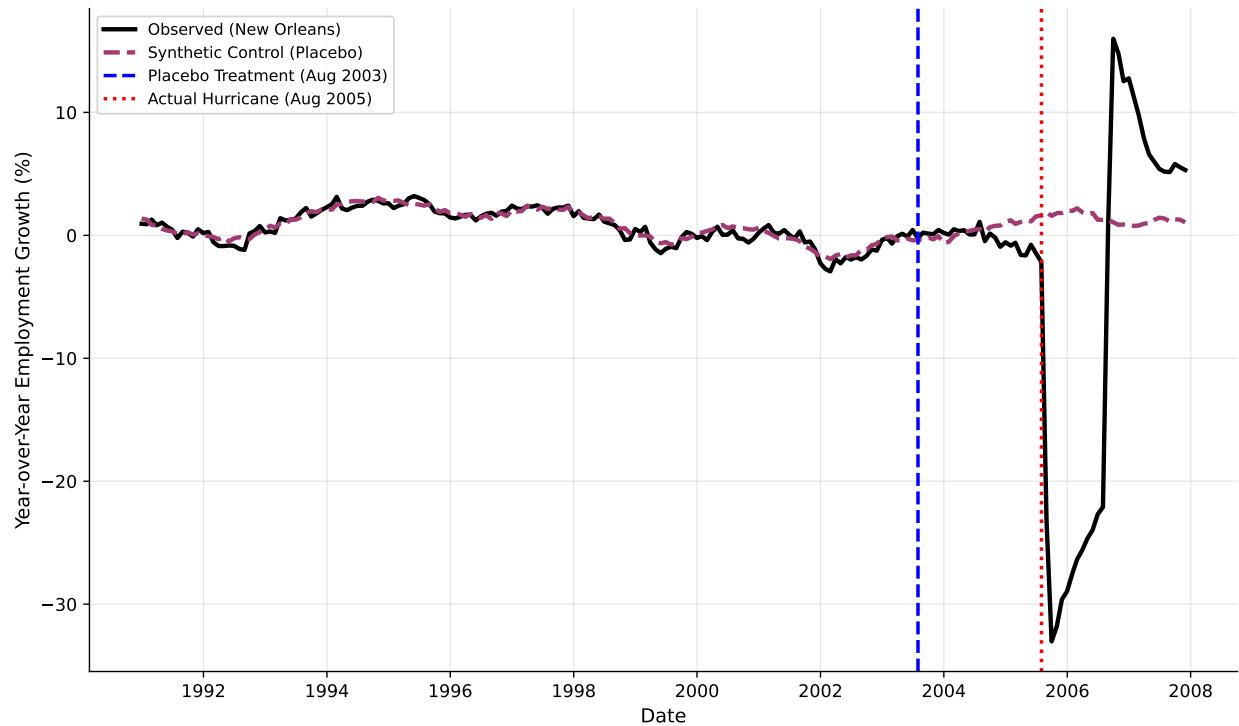


Figure 3: In-Time Placebo Test

Note: Validation test with placebo treatment placed in 2003 (dashed vertical line). Gap shows no systematic collapse at placebo timing, with catastrophic impact appearing only at actual hurricane time in 2005.

Appendix

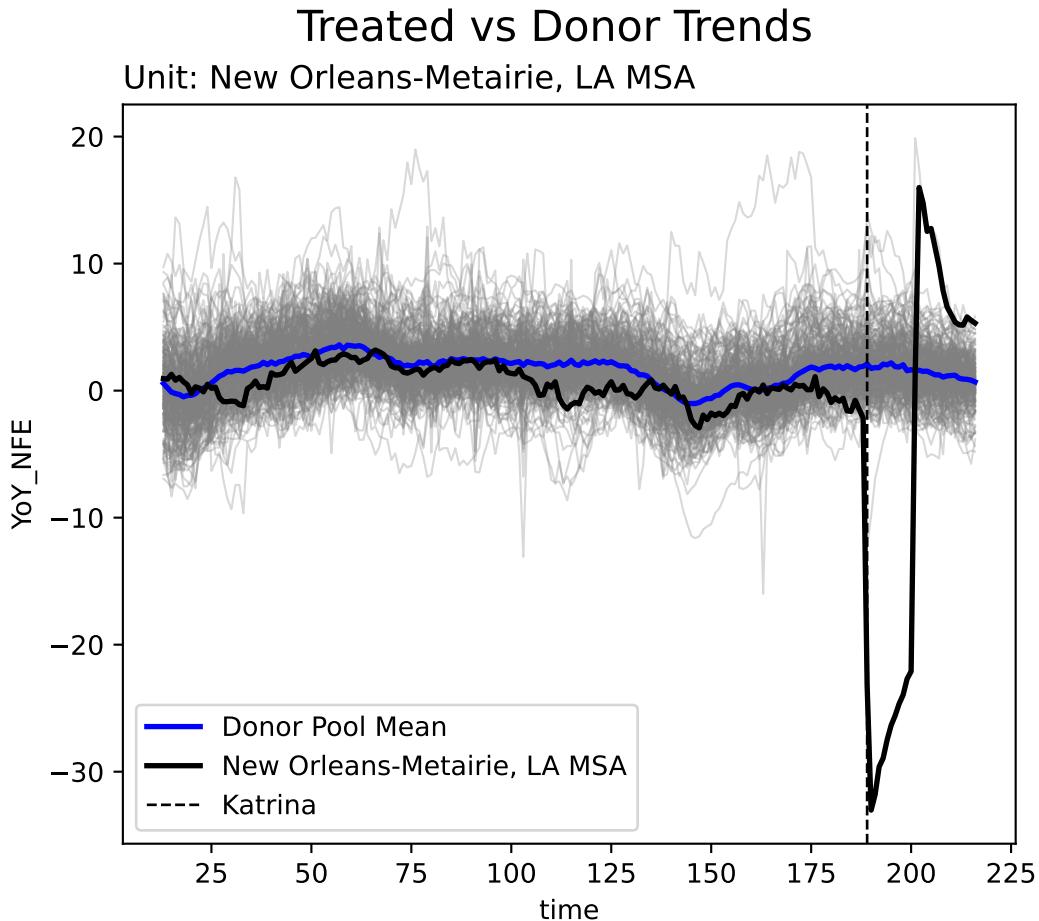


Figure A1: Naive Comparison - New Orleans vs Donor Pool Average

Note: New Orleans (black line) versus simple average of 273 donor MSAs (blue line). Gray lines show individual donor pool MSAs. Vertical dashed line indicates Hurricane Katrina (August 2005).

Note: Gap between actual New Orleans employment growth and synthetic control during pre-treatment period (1990-2005). Gap fluctuates tightly around zero throughout 15-year baseline, validating the identification assumption.

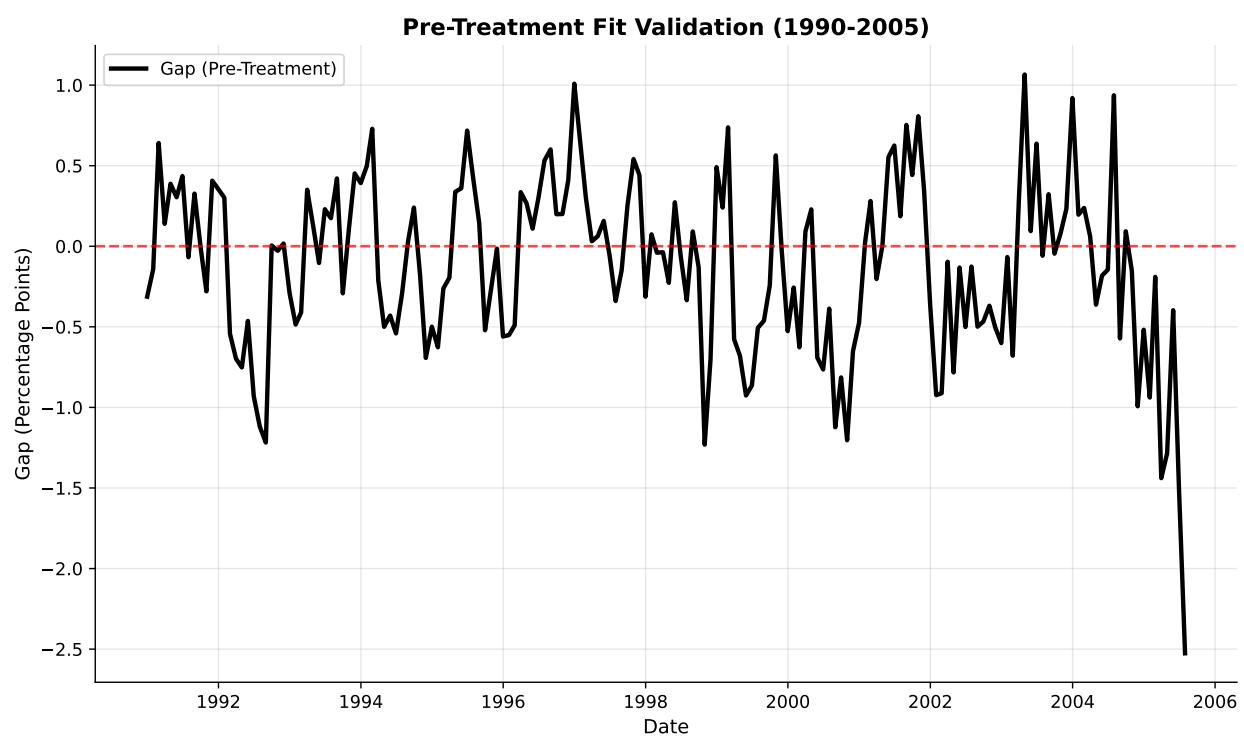


Figure A2: Pre-Treatment Fit Validation